

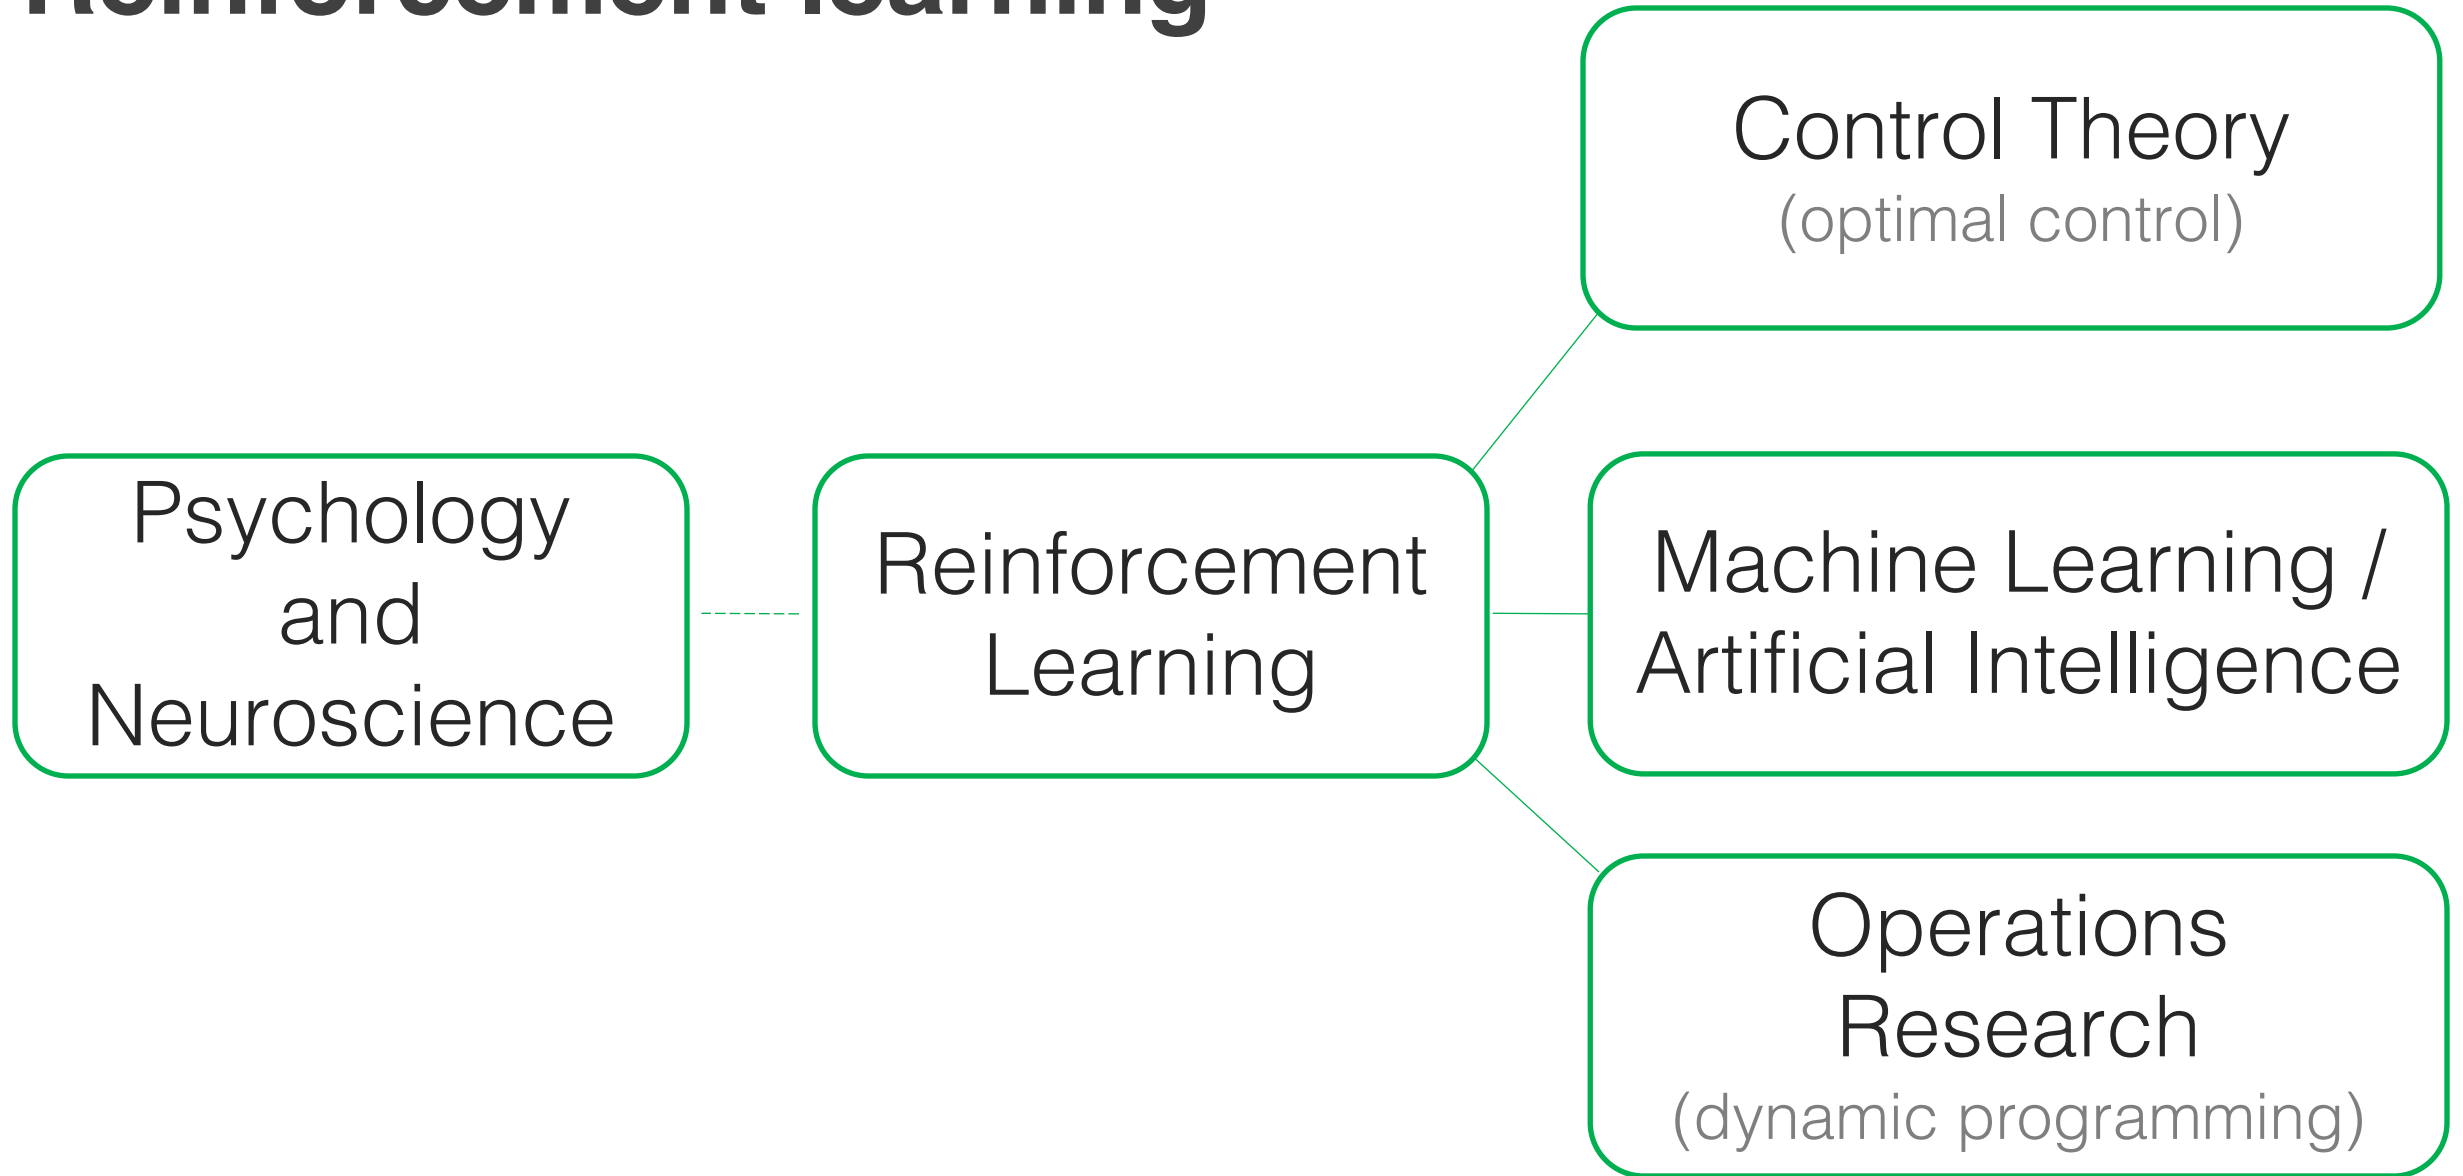
Reinforcement Learning I

Lecture 19

Types of machine learning

	Supervised Learning	Unsupervised Learning	Reinforcement Learning
Goal	Predict ...from examples	Describe ...structure in data	Strategize learn by trial and error
Data	(x, y)	x	delayed feedback
Types	<ul style="list-style-type: none">• Classification• Regression	<ul style="list-style-type: none">• Density estimation• Clustering• Dimensionality reduction• Anomaly detection	<ul style="list-style-type: none">• Model-free learning• Model-based learning

Reinforcement learning



Resources

This reinforcement learning series draws heavily on these resources

Sutton and Barto, 1998

Reinforcement Learning: An Introduction

Draft of updated edition available free online:

<http://www.incompleteideas.net/book/the-book-2nd.html>




Image from Amazon.com (where the book may be purchased)

David Silver, 2015

University College London

Advanced Topics 2015 (COMPM050/COMPGI13)

Course website:

<http://www0.cs.ucl.ac.uk/staff/D.Silver/web/Teaching.html>

Video series:

<https://www.youtube.com/watch?v=2pWv7GOvuf0&list=PL7-jPKtc4r78-wCZcQn5lqyuWhBZ8fOxT>

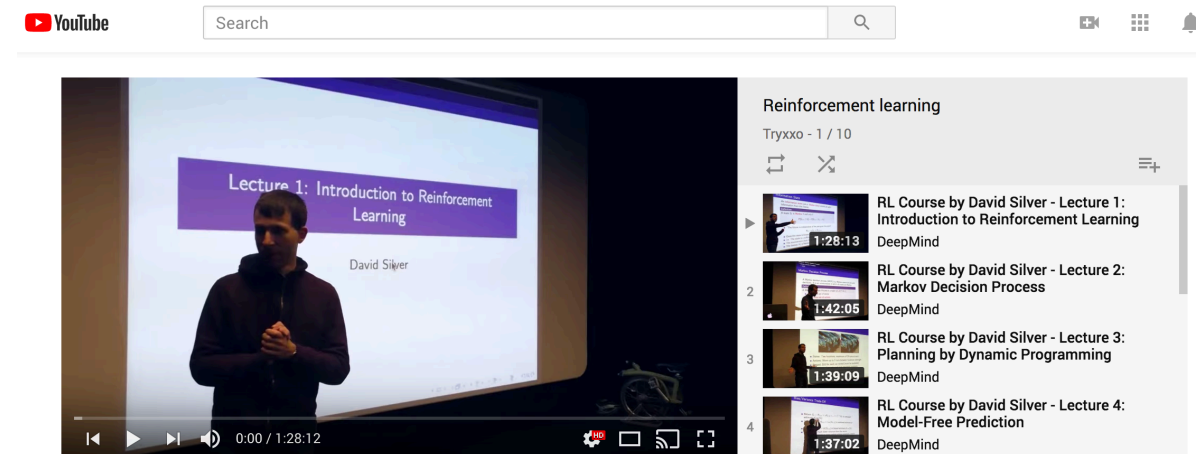


Image from Youtube.com

Reinforcement Learning

Goal: select actions to maximize total long-term rewards

Sequential decision making

Challenge: an action needs to be taken at each step

Evaluation of rewards versus instruction (examples of correct actions)

Challenge: this leads to a trial-and-error approach to learning

May be better to sacrifice immediate reward for long-term gains

Challenge: exploration (of untried actions) vs exploitation (of current knowledge)

Rewards may be delayed

Challenge: credit assignment: which action(s) led to the reward(s)?

Reinforcement Learning Examples

Winning at Atari: <https://youtu.be/V1eYniJ0Rnk>

Balancing an inverted pendulum: https://youtu.be/b1c0N_Fs9wc

Flipping pancakes: https://youtu.be/W_gxLKSsSIE

Car Drifting: <https://youtu.be/opsmd5yuBF0>


RL is a unifying framework for a wide range of problems

Multi-armed Bandit



You walk into a casino...


Slot Machine (a.k.a. one-armed bandit)



win or lose

Define this as reward is “1” for win or “0” for lose

Multi-armed bandit problem



Trial/episode: play one machine

Action: pick one machine to play (one action per trial/episode)

Reward: how much you win or lose

- Each machine has an unknown probability of payoff/reward
- The rewards are stochastic (their distributions are unknown)

Action-Value: expected reward for taking each action

State: only 1 “state” in this problem - our environment doesn’t change

Policy: create a policy How do we choose actions to maximize our total rewards?

- If we knew the best machine, we’d always pick it
- This is what we want to learn

Multi-armed Bandit Demo

<http://dev.kylebradbury.org/bandit/>

Multi-armed bandit

The “true” **action-value** of an action is $q_*(a)$

Our estimated **action-value** at the t^{th} play is $q_t(a)$

If action a has been chosen k_a times prior to t :

$$q_t(a) = \frac{r_1 + r_2 + \cdots + r_{k_a}}{k_a}$$

As we take action a more, our action-value estimates improve

Multi-armed bandit policies, $\pi(s)$

Greedy action:

Select $a^* = \arg \max_a q_t(a)$

Problem: if the initial rewards are not representative, this will be suboptimal

ϵ -Greedy methods:

Select a^* with probability $1 - \epsilon$, otherwise, randomly select another option

Problem: in the long run, this will waste reward once the best action is known


Solution: reduce ϵ over time

Alternative:



Select the action probabilities based on the expected value

Probability of selecting action $P(a) = \frac{\exp(q_t(a))}{\sum_{b=1}^n \exp(q_t(b))}$

10-Armed Bandit Example



Note: Each distribution has a mean $q_*(a)$ with unit variance



Sutton and Barto, 2018

Next steps

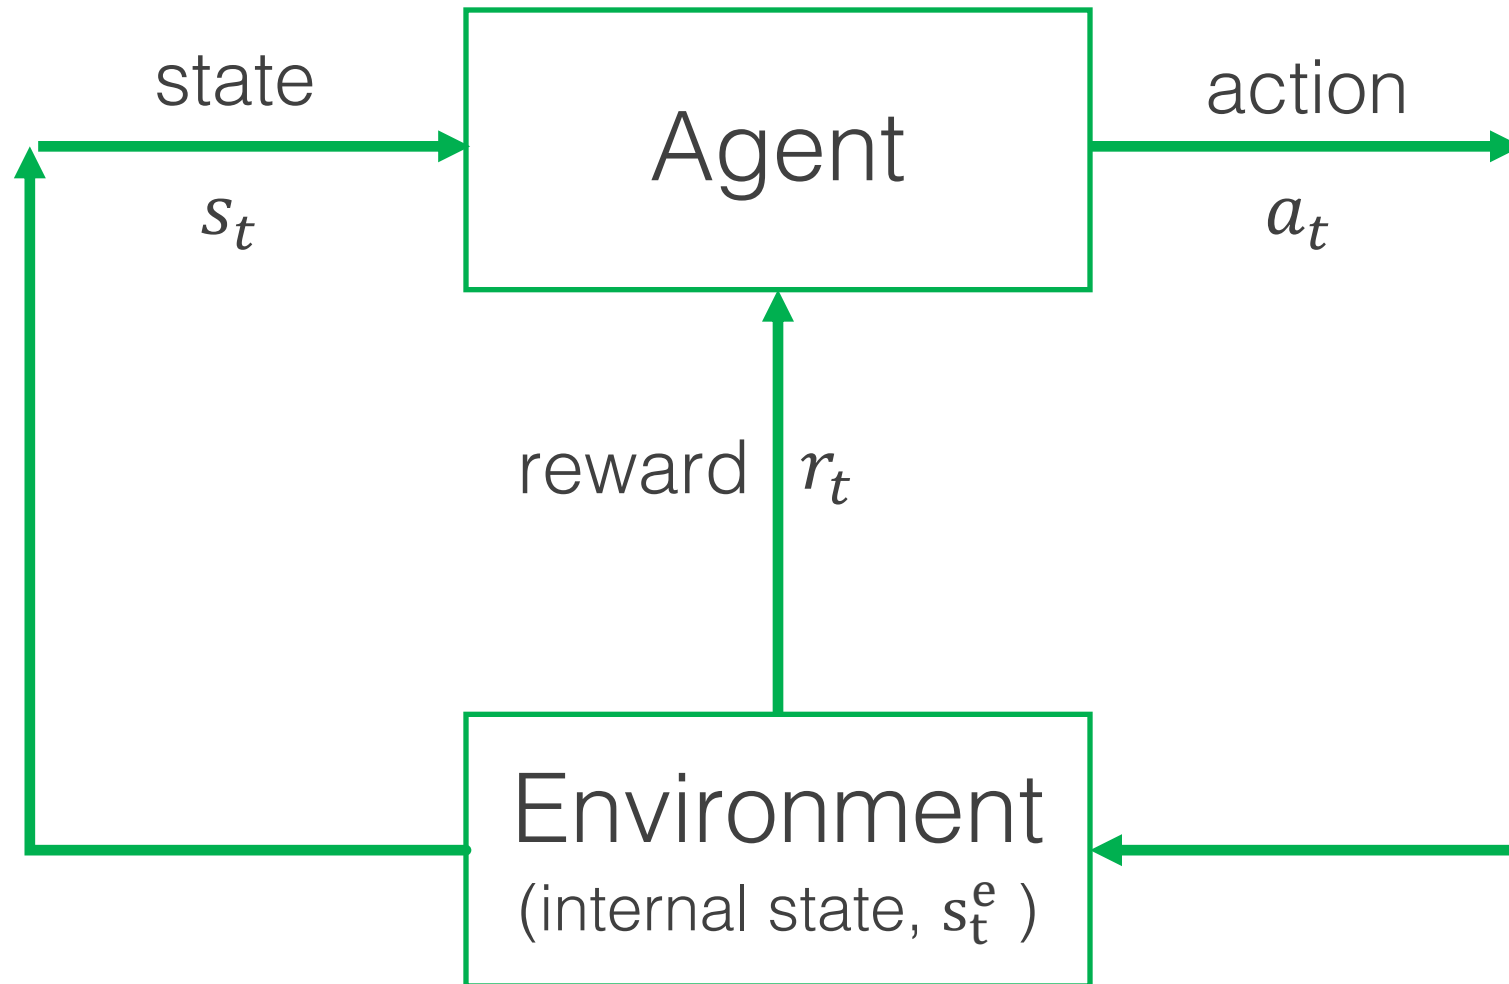
The multi-armed bandit only has 1 state, but the full RL problem learns policies when there are many states

State representations and Markov decision processes (MDPs)
(with an aside on Markov processes)

Mathematically formulating the RL problem with MDPs

Methods for solving RL problems in practice

Agent-environment Interaction



Agent at each step t ...

- Executes action a_t
- Receives state, s_t
- Receives scalar reward, r_t

Environment at each step t ...

- Receives action a_t
- Emits state, s_{t+1}
- Emits scalar reward, r_{t+1}


Actions: choices made by the agent

States: basis on which choices are made

Rewards: define the agent's goals

David Silver, 2015

Reinforcement Learning Components




Policy (agent behavior), $\pi(s)$

Reward function (the goal), r_t

Value functions (expected returns), $v(s), q(s)$

Maze Example: Policy, Value, and Reward



Each location in the maze represents a **state**

The **reward** is -1 for each step the agent is in the maze

Available **actions**: move $\uparrow, \downarrow, \leftarrow, \rightarrow$ (as long as that path is not blocked)

Adapted from David Silver, 2015

Policy $\pi(s)$

(which actions to take in each state)

Start			
	↓		
→	→	↓	←
↑		↓	
	→	→	↓
	↑		↓
→	↑		↓
Exit			

Reward r_t

(rewards are received as you transition OUT OF the state)

Start			
	-1		
-1	-1	-1	-1
-1		-1	
	-1	-1	-1
	-1		-1
-1	-1		-1
Exit			


State Value $v_\pi(s)$

(expected cumulative rewards starting from current state **if** we follow the policy)

Start			
	-8		
-8	-7	-6	-7
-9		-5	
	-5	-4	-3
	-6		-2
-8	-7		-1
Exit			

Adapted from David Silver, 2015

Policy



Policy, $\pi(s)$

- Agent's way of behaving at a given time
- Maps state to actions

Deterministic: $a = \pi(s)$

Stochastic: $\pi(a|s) = P(a_t = a | s_t = s)$

Helps us “explore” the state space

RL tries to learn the “best” policy

Goals and rewards

Rewards are the **only way** of communicating what to accomplish

Ex 1: Robot learning a maze

- 0 until it escapes, then +1 when it does
- -1 until it escapes (encourages it to escape quickly)

Ex 2: Robot collecting empty soda cans

- +1 for each empty soda can
- Negative rewards for bumping into things

Chess: what if we set +1 for capturing a piece?
(it may not win the game and still maximize rewards)

What you want achieved not **how**

Returns / cumulative reward

Episodic tasks (finite number, T , of steps, then reset)

$$G_t = r_{t+1} + r_{t+2} + \dots + r_T$$


Continuing tasks with discounting ($T \rightarrow \infty$)

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

where $0 \leq \gamma \leq 1$ is the discount rate

This makes the agent care more about immediate rewards

Value functions



State Value function, $v_\pi(s)$

- How “good” is it to be in a state, s_t then follow policy π to choose actions
- Total expected rewards

$$v_\pi(s) = E_\pi[G_t | s_t = s]$$


Action Value function, $q_\pi(s, a)$

- How “good” is it to be in a state, s , take action a , then follow policy π to choose actions
- Total expected rewards

$$q_\pi(s, a) = E_\pi[G_t | s_t = s, a_t = a]$$

Where
$$G_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

Model



Model

Transitions: predicts what state the environment will transition to next

$$P_{ss'}^a = P(s_{t+1} = s' | s_t = s, a_t = a)$$

Rewards: predicts the next reward given an action


$$R_s^a = E[r_{t+1} | s_t = s, a_t = a]$$

“Planning” is the process of using these predictions

Model-based RL uses a model

Model-free RL does not use a model

Reinforcement Learning Components



Policy (agent behavior), $\pi(s)$

- Determines action given current state
- Agent's way of behaving at a given time

Reward function (the goal), r_t

- Maps state of the environment to a reward that describes the state desirability
- Objective is to **maximize total rewards**

Value (expected returns), $v(s, a), q(s)$

- Total expected reward from a state
- How "good" is each state